

■ Full List of Fine-Tuning Frameworks & Tools (2025)

1. Hugging Face Transformers

Most popular for NLP/LLM fine-tuning. Integrates with Trainer, PEFT, Datasets, etc.
<https://huggingface.co/docs/transformers>

2. PEFT

Techniques: LoRA, Prefix-Tuning, Prompt-Tuning. Lightweight tuning of large models.
<https://github.com/huggingface/peft>

3. Unsloth

Fast, memory-efficient LoRA fine-tuning. Fine-tunes 7B models on 12GB GPUs.
<https://github.com/unslothai/unsloth>

4. Axolotl

Complete fine-tuning framework for open LLMs. Supports LoRA, QLoRA, DPO, etc.
<https://github.com/OpenAccess-AI-Collective/axolotl>

5. DeepSpeed

Distributed fine-tuning & training. Advanced memory optimization (ZeRO, Offloading).
<https://www.deepspeed.ai>

6. QLoRA

Quantized Low-Rank Adaptation. Fine-tune 33B/65B+ models on limited GPUs.
<https://github.com/artidoro/qlora>

7. LoRA

Efficient fine-tuning method. Plugged into PEFT, DeepSpeed, Axolotl.
<https://github.com/microsoft/LoRA>

8. TRL / TRLLM

Reinforcement Learning fine-tuning (RLHF, PPO, DPO). <https://github.com/huggingface/trl>

9. Colossal-AI

Distributed training for large models. <https://www.colossalai.org>

10. Lit-GPT

Fast GPT fine-tuning using Lightning AI. <https://github.com/Lightning-AI/lit-gpt>

11. FastChat

Fine-tune & serve chat models. <https://github.com/lm-sys/FastChat>

12. OpenLLM

Production-grade LLM serving & tuning. <https://github.com/bentoml/OpenLLM>

13. Text Generation Inference (TGI)

Inference server for LLMs. <https://github.com/huggingface/text-generation-inference>

14. LLaMA Factory

Easy fine-tuning for LLaMA models. <https://github.com/hiyouga/LLaMA-Factory>

15. S-LoRA / SpQR / EXL2 / GPTQ

Advanced quantization and sparse tuning methods. <https://github.com/casper-hansen/AutoGPTQ>